

Deep Reinforcement Learning Practical Session

Large Scale Machine Learning - Mines ParisTech

March 2022

1 Introduction

All of the instructions are on the [notebook](#). We recommend staying in Colab. You should follow the notebook instructions for the deliverables. Here you will find some theoretical background for each question.

Some definitions:

- A *trajectory* τ is a sequence $(s_1, a_1, s_2, a_2 \dots s_T, a_T)$ of the visited states and actions. The probability of a trajectory under the policy π is $p_\pi(\tau) = \pi(\tau) = p(s_1) \prod_{t=1}^T \pi(a_t|s_t)p(s_{t+1}|s_t, a_t)$.
- The discounted cumulative reward is $R_t = \sum_{k=0}^T \gamma^k r(s_{t+k}, a_{t+k})$. Our goal is to minimize $J(\pi) = \mathbb{E}_{\tau \sim p_\pi(\tau)} [R_1]$.
- The *Q-function* $Q^\pi(s_t, a_t) = \mathbb{E}_{p_\pi} [R_t | s_t, a_t] = \sum_{t'=t}^T \mathbb{E}_{p_\pi} [\gamma^{t'-t} r(s_{t'}, a_{t'}) | s_t, a_t]$ is the reward-to-go from state s_t if we pick the action a_t and then follow π .
- The *value function* $V^\pi(s_t) = \mathbb{E}_{p_\pi} [R_t | s_t] = \sum_{t'=t}^T \mathbb{E}_{p_\pi} [\gamma^{t'-t} r(s_{t'}, a_{t'}) | s_t] = \mathbb{E}_{a_t \sim \pi(a_t | s_t)} [Q^\pi(s_t, a_t)]$ is the reward-to-go from state s_t if we follow the policy π .
- The *advantage function* $A^\pi(s_t, a_t) = Q^\pi(s_t, a_t) - V^\pi(s_t)$ represents the value of a particular action a_t with respect to the average action taken by π at state s_t .

2 Policy gradient

The idea is to do stochastic gradient ascent on the objective $J(\pi) = \mathbb{E}_{\tau \sim p_\pi(\tau)} [R_1]$ to find the best parameterized stochastic policy π_θ . A simple example is the Reinforce algorithm [4].


- 
1. sample $\{\tau^i\}$ from $\pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$ (run the policy)
 2. $\nabla_\theta J(\theta) \approx \sum_i (\sum_t \nabla_\theta \log \pi_\theta(\mathbf{a}_t^i | \mathbf{s}_t^i)) (\sum_t r(\mathbf{s}_t^i, \mathbf{a}_t^i))$
 3. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Figure 1: Reinforce algorithm (source: [1])

Reinforce uses the *likelihood ratio trick* to get rid of the terms depending on the dynamics:

$$\nabla_\theta J(\theta) = \nabla_\theta \int_\tau p_\theta(\tau) r(\tau) d\tau = \int_\tau p_\theta(\tau) \nabla_\theta \log p_\theta(\tau) r(\tau) d\tau = \mathbb{E}_{\tau \sim p_\theta(\tau)} \left[\nabla_\theta \sum_t \log \pi_\theta(a_t | s_t) r(\tau) \right] \quad (1)$$

Step 2 of the algorithm uses samples to approximate the expectation. Step 3 is gradient ascent, making good sampled trajectories more likely, and bad trajectories less likely.

2.1 CartPole - Compare different hyper-parameters

Run some experiments in the discrete [cartpole](#) environment. Two interesting parameters to tune are:

reward.to.go. We usually replace $\sum_{t=1}^T \gamma^{t-1} r(s_t^i, a_t^i)$ with $\sum_{t'=t}^T \gamma^{t'-t} r(s_{t'}^i, a_{t'}^i)$, i.e. we hold policies accountable for future rewards only, which reduces the variance but might bring a bias when $\gamma < 1$.

batch.size. Number of state-action pairs sampled while acting according to the current policy at each iteration. A bigger batch size should also reduce the variance.

2.2 LunarLander - Use a neural network baseline

Try using a neural network baseline in the more demanding continuous [lunar lander](#) environment.

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \left(\sum_{t'=t}^T \gamma^{t'-t} r(s_{t'}^i, a_{t'}^i) - b \right) \quad (2)$$

Reinforce uses the MC (Monte Carlo) estimator of the Q-function with a single sample $\hat{Q}^{\pi}(s_t, a_t) = \sum_{t'=t}^T \gamma^{t'-t} r(s_{t'}^i, a_{t'}^i)$, hence the high variance. To reduce the variance, we can add a state-dependant baseline (fancy term meaning some quantity to compare against) that approximates the value function $b = \hat{V}^{\pi}(s_t)$, so that we compare the obtained returns with the expected returns. This choice of baseline results in an approximation of the advantage function $\hat{A}^{\pi}(s_t, a_t)$.

The goal is to approximate the value function with a neural network V_{ϕ} . To do so, we use the return-to-go MC targets $y_i = \sum_{t'} \gamma^{t'-t} r(s_{t'}^i, a_{t'}^i)$ and minimize the supervised regression loss $\frac{1}{2} \sum_i \|V_{\phi}^{\pi}(s_t^i) - y_i\|^2$.

- Remark: In practice, we compute normalized targets because it is easier for the network to predict returns with zero mean and standard deviation of one.
- Remark: The PyTorch function [detach\(\)](#) stops the gradients from propagating back to the value function network when updating the policy.

3 Q-learning

The goal is to recover the optimal Q-function Q^* , which we know verifies the optimal Bellman equation (derived from dynamic programming). A similar equation holds for the value function, but the *max* operator lies outside of the expectation, which prevents us from using samples to approximate the expectation. One advantage of approximating Q^* rather than Q^{π} is that the algorithms become *off-policy*: theoretically, the training data can come from any policy, as opposed to *on-policy* algorithms such as policy gradient which require data coming from π to approximate Q^{π} .

$$Q^*(s_t, a_t) = \mathbb{E}_{s_{t+1} \sim p(\cdot | s_t, a_t)} \left[r(s_t, a_t) + \gamma \max_{a_{t+1}} Q^*(s_{t+1}, a_{t+1}) \right] \quad (3)$$

To do so, we minimize the TD (temporal difference) error $\sum_i \|Q_{\phi}(s_i, a_i) - y_i\|^2$ where $y_i = r(s_i, a_i) + \gamma \max_{a'_i} Q_{\phi}(s'_i, a'_i)$. A very popular Q-learning algorithm is DQN [3]. Let's cover the two main ideas:

- *Replay buffer.* As we interact with the environment sequentially, the samples are strongly correlated which can lead to local over-fitting. The idea is to have a buffer \mathcal{B} where we add the collected data, and then update the Q-function from data sampled uniformly from \mathcal{B} .
- *Target network.* We do not regress the Q-function to the TD targets y until convergence because the optimal estimate from the available data will be very poor in early stages. This means that the targets

y are going to change very frequently, which doesn't help with stability. The idea is to do regular backups of the parameters ϕ into a new target network $Q_{\phi'}$ used to compute these targets.

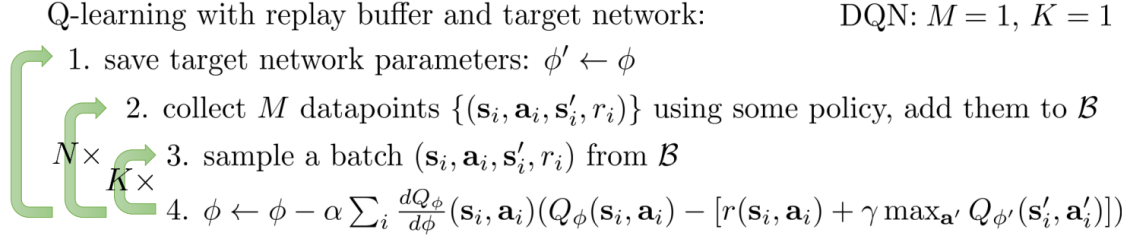


Figure 2: DQN algorithm (source: [1])

3.1 MountainCar - Compare different rewards

Solving the [mountain car](#) task is very difficult with the original reward, which gives -1 at every timestep until the car reaches the flag. In order to make any progress, the car needs to reach the flag first, and relies only on exploration to do so. A better strategy would be to reward an increase in mechanical energy (potential plus kinetic) at every time step. Try out this alternative reward in the notebook.

3.2 LunarLander - Implement double DQN

One problem with DQN is overestimation when computing the targets. Since $Q_{\phi'}$ is not perfect, $\max_{\mathbf{a}'} Q_{\phi'}(s', a')$ will be an overestimation if there is any upwards noise. If we could select the action according to a different network with decorrelated noise the problem would go away. Since we already have two networks, the idea is to use $y = r + \gamma Q_{\phi'}(s', \arg \max_{\mathbf{a}'} Q_\phi(s', a'))$ rather than $y = r + \gamma Q_{\phi'}(s', \arg \max_{\mathbf{a}'} Q_{\phi'}(s', a'))$.

You can compare the stability of the learning process with or without DDQN in the lunar lander environment.

4 Actor Critic

Actor-critic algorithms implement a network for both the policy (actor) and the value or Q-function (critic). DDPG [2] is a very popular example. It is an interesting algorithm because it was introduced as an alternative to policy gradient for deterministic policies, and can also be seen as an alternative to Q-learning for continuous actions. A key feature is that it works off-policy.

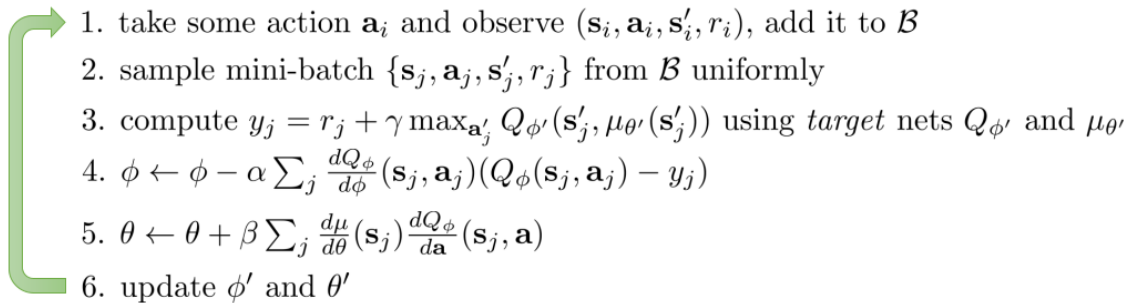


Figure 3: DDPG: deep deterministic policy gradient (source: [1])

4.1 Pendulum - Implement the critic and actor updates

You can try out your code in the [pendulum](#) environment.

- Hint: For the actor, use the chain rule: $\nabla_{\theta} Q_{\phi}(s, \mu_{\theta}(s)) = \nabla_a Q_{\phi}(s, a = \mu_{\theta}(s)) \nabla_{\theta} \mu_{\theta}(s)$.

References

- [1] Sergey Levine. *CS285: Deep Reinforcement Learning*. Fa2019. URL: <http://rail.eecs.berkeley.edu/deeprlcourse/>.
- [2] Timothy P Lillicrap et al. “Continuous control with deep reinforcement learning”. In: *arXiv preprint arXiv:1509.02971* (2015).
- [3] Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (2015), pp. 529–533.
- [4] Ronald J Williams. “Simple statistical gradient-following algorithms for connectionist reinforcement learning”. In: *Machine learning* 8.3-4 (1992), pp. 229–256.